

# Worksheet 5 output

## Task 1: Data Understanding and Visualization:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
from PIL import Image
import os
import glob
import numpy as np
import matplotlib.pyplot as plt
import random
```

```
train_path = "/content/drive/MyDrive/Artificial Intelligence and machine learning/FruitinAmazon/FruitinAmazon/train"
test_path = "/content/drive/MyDrive/Artificial Intelligence and machine learning/FruitinAmazon/FruitinAmazon/test"
```

```
os.listdir(train_path)
```

```
['tucuma', 'guarana', 'cupuacu', 'graviola', 'acai', 'pupunha']
```

```
visualise(train_path)
```



```
        img.verify()
    except Exception as e:
        print(f"Corrupted image found {image}")
    print("No corrupted image found")

check_for_corrupted(train_path)
```

No corrupted image found

[+ Code](#)[+ Text](#)

## Task 2: Loading and Preprocessing Image Data in keras:

[+ Code](#)[+ Text](#)

```
img_height = 128
img_width = 128
batch_size = 32
validation_split = 0.2

rescale = tf.keras.layers.Rescaling(1./255)

train_ds = tf.keras.utils.image_dataset_from_directory(
    train_path,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=True,
    validation_split=validation_split,
    subset='training',
    seed=123
```

```
test_ds = test_ds.map(lambda x, y: (rescale(x), y))
```

Found 90 files belonging to 6 classes.  
Using 72 files for training.  
Found 30 files belonging to 6 classes.  
Found 30 files belonging to 6 classes.

[+ Code](#)[+ Text](#)

## Task 3 - Implement a CNN with Convolutional Architecture:

```
model = keras.Sequential([
    layers.Conv2D(32, (3, 3), padding='same', strides=1, activation='relu', input_shape=(128, 128, 3)),
    layers.MaxPooling2D((2, 2), strides=2),
    layers.Conv2D(32, (3, 3), padding='same', strides=1, activation='relu'),
    layers.MaxPooling2D((2, 2), strides=2),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(128, activation='relu'),
    layers.Dense(6, activation="softmax")
])
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```
model.summary()
```

model.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_1 (Conv2D)	(None, 64, 64, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 32)	0
flatten (Flatten)	(None, 32768)	0
dense (Dense)	(None, 64)	2,097,216
dense_1 (Dense)	(None, 128)	8,320
dense_2 (Dense)	(None, 6)	774

Total params: 2,116,454 (8.07 MB)  
Trainable params: 2,116,454 (8.07 MB)  
Non-trainable params: 0 (0.00 B)

## Task 4: Compile the Model

model.compile(  
 optimizer='adam',  
 loss='sparse\_categorical\_crossentropy',  
 metrics=['accuracy']  
)

batch\_size=16,  
validation\_data=val\_ds,  
callbacks=callbacks  
)

Epoch 1/250  
3/3 15s 6s/step - accuracy: 0.1623 - loss: 1.8829 - val\_accuracy: 0.3333 - val\_loss: 1.7234  
Epoch 2/250  
3/3 3s 821ms/step - accuracy: 0.4062 - loss: 1.6800 - val\_accuracy: 0.3333 - val\_loss: 1.6413  
Epoch 3/250  
3/3 4s 604ms/step - accuracy: 0.3958 - loss: 1.5225 - val\_accuracy: 0.1667 - val\_loss: 1.7050  
Epoch 4/250  
3/3 3s 996ms/step - accuracy: 0.3832 - loss: 1.3420 - val\_accuracy: 0.5333 - val\_loss: 1.3762  
Epoch 5/250  
3/3 4s 1s/step - accuracy: 0.6337 - loss: 1.0028 - val\_accuracy: 0.3333 - val\_loss: 1.3530  
Epoch 6/250  
3/3 4s 649ms/step - accuracy: 0.6771 - loss: 0.8534 - val\_accuracy: 0.4000 - val\_loss: 1.4508  
Epoch 7/250  
3/3 2s 619ms/step - accuracy: 0.7422 - loss: 0.7113 - val\_accuracy: 0.4000 - val\_loss: 1.3276  
Epoch 8/250  
3/3 2s 701ms/step - accuracy: 0.8906 - loss: 0.5053 - val\_accuracy: 0.6333 - val\_loss: 1.0230  
Epoch 9/250  
3/3 3s 952ms/step - accuracy: 0.9488 - loss: 0.2670 - val\_accuracy: 0.7000 - val\_loss: 0.9795  
Epoch 10/250  
3/3 4s 667ms/step - accuracy: 0.9705 - loss: 0.1708 - val\_accuracy: 0.6333 - val\_loss: 1.3013  
Epoch 11/250  
3/3 2s 630ms/step - accuracy: 0.9410 - loss: 0.1909 - val\_accuracy: 0.6667 - val\_loss: 1.0288  
Epoch 12/250  
3/3 2s 666ms/step - accuracy: 0.9518 - loss: 0.1235 - val\_accuracy: 0.5667 - val\_loss: 1.2649  
Epoch 13/250  
3/3 2s 702ms/step - accuracy: 0.9891 - loss: 0.0734 - val\_accuracy: 0.5333 - val\_loss: 1.2460

```
plt.show()
```



## Task 5: Evaluate the Model

[+ Code](#)[+ Text](#)

```
test_loss, test_acc = model.evaluate(test_ds)
print(f"Test Accuracy: {test_acc:.4f}")
```

```
1/1 ————— 0s 329ms/step - accuracy: 0.7000 - loss: 0.9795
Test Accuracy: 0.7000
```

```
model.save('bestest_model.h5')
```

```
bestest_model = tf.keras.models.load_model('bestest_model.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via "model.save()" or "keras.saving.save\_model(model)". This file format is considered legacy. We recommend using instead the nat

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. "model.compile\_metrics" will be empty until you train or evaluate the model.

## Task 7: Predictions and Classification Report

```
import numpy as np
from sklearn.metrics import classification_report
```

```
y_pred_prob = bestest_model.predict(test_ds)
y_pred = np.argmax(y_pred_prob, axis=1)
```

```
y_true = []
for images, labels in test_ds.unbatch():
```

```
    y_true.append(labels)
```

```
y_true = np.array(y_true)
print(y_true)
print(classification_report(y_true, y_pred))
```

```
1/1 ————— 1s 689ms/step
[0 0 0 0 1 1 1 1 1 2 2 2 2 3 3 3 3 3 4 4 4 4 4 4 5 5 5 5]
      precision    recall  f1-score   support

     0       0.67       0.80       0.73         5
     1       0.50       0.60       0.55         5
     2       0.71       1.00       0.83         5
     3       0.67       0.80       0.73         5
     4       1.00       0.60       0.75         5
     5       1.00       0.40       0.57         5

 accuracy          0.70         0.70         0.70        30
 macro avg          0.76         0.70         0.69        30
 weighted avg       0.76         0.70         0.69        30
```